

# CSE 405: Machine Learning

## Supervised Learning

**Dr Muhammad Abul Hasan**



Department of Computer Science and Engineering  
Green University of Bangladesh  
`muhammad.hasan@cse.green.edu.bd`

Fall 2023



**“** *I am always doing that which I cannot do, in order that I may learn how to do it.*

*– Pablo Picasso*

**”**

# Supervised Learning



# What is Supervised Learning?

## Definition

Supervised learning is a type of machine learning where the algorithm learns from labeled training data to make predictions or decisions without being explicitly programmed.

- Labeled training data
- Predictions or decisions
- Example: Email spam classification



# The Supervised Learning Workflow

## Steps in Supervised Learning

- 1 Data collection
- 2 Data preprocessing
- 3 Model selection
- 4 Model training
- 5 Model evaluation
- 6 Prediction



# Types of Supervised Learning Problems

## Regression

Predicting a continuous output variable (e.g., price, temperature).

## Classification

Assigning data to predefined classes or categories (e.g., spam or not spam, image recognition).



## Learning a Class from Examples

- Class C of a “family car”
  - ▣ Prediction: Is car  $x$  a family car?
  - ▣ Knowledge extraction: What do people expect from a family car?
- Output: Positive (+) and negative (−) examples
- Input representation:
  - ▣  $x_1$ : price
  - ▣  $x_2$ : engine power





## Training set $X$


A training set is expressed as:

$$X = \{\mathbf{x}^t, r^t\}_{t=1}^N$$

where,

$$\mathbf{x}^t = \begin{bmatrix} x_1^t \\ x_2^t \end{bmatrix},$$


$$r^t = \begin{cases} 1 & \text{if } \mathbf{x}^t \text{ is positive} \\ 0 & \text{if } \mathbf{x}^t \text{ is negative} \end{cases}$$



# Linear Regression




## Problem Setup



- Want to predict a scalar  $t$  as a function of a scalar  $x$
- Given a dataset of pairs  $(\mathbf{x}^{(i)}, t^{(i)})_{i=1}^N$
- The  $\mathbf{x}^{(i)}$  are called inputs, and the  $t^{(i)}$  are called targets.



## Problem Setup



- Model:  $y$  is a linear function of  $x$ :

$$y = wx + b$$

- $y$  is the prediction
- $w$  is the weight
- $b$  is the bias
- $w$  and  $b$  together are the parameters
- Settings of the parameters are called hypotheses



## Problem Setup

- Loss function: squared error (says how bad the fit is)



$$\mathcal{L}(y, t) = \frac{1}{2}(y - t)^2$$

- $y - t$  is the residual, and we want to make this small in magnitude
- The  $\frac{1}{2}$  factor is just to make the calculations convenient.
- Cost function: loss function averaged over all training examples

$$\begin{aligned}\mathcal{J}(w, b) &= \frac{1}{2N} \sum_{i=1}^N (y^{(i)} - t^{(i)})^2 \\ &= \frac{1}{2N} \sum_{i=1}^N (wx^{(i)} + b - t^{(i)})^2\end{aligned}$$



## Problem Setup



## Problem Setup

- Suppose we have multiple inputs  $x_1, \dots, x_D$ . This is referred to as multivariable regression.
- This is no different than the single input case, just harder to visualize.
- Linear model:

$$y = \sum_j w_j x_j + b$$



## How to decide the value of $w_j$ and $b$ ?

Gradient descent. It is an optimization algorithm.





## Solving the optimization problem

- We defined a cost function. This is what we'd like to minimize.
- Recall from calculus class: minimum of a smooth function (if it exists) occurs at a critical point, i.e. point where the derivative is zero.
- Multivariate generalization: set the partial derivatives to zero. We call this a direct solution.



## Gradient Descent

- Gradient descent is an iterative algorithm, which means we apply an update repeatedly until some criterion is met.
- We initialize the weights to something reasonable (e.g. all zeros) and repeatedly adjust them in the direction of the steepest descent.



## Gradient Descent

### ■ Observe:

- ▣ if  $\frac{\partial \mathcal{J}}{\partial w_j} > 0$ , then increasing  $w_j$  increases  $\mathcal{J}$ .
- ▣ if  $\frac{\partial \mathcal{J}}{\partial w_j} < 0$ , then increasing  $w_j$  decreases  $\mathcal{J}$ .

### ■ The following update decreases the cost function:

$$\begin{aligned} w_j &\leftarrow w_j - \alpha \frac{\partial \mathcal{J}}{\partial w_j} \\ &= w_j - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)} \end{aligned}$$

### ■ $\alpha$ is a learning rate. The larger it is, the faster $\mathbf{w}$ changes.

- ▣ We will see later how to tune the learning rate, but values are typically small, e.g. 0.01 or 0.0001.



## Gradient Descent

- This gets its name from the gradient:

$$\frac{\partial \mathcal{J}}{\partial \mathbf{w}} = \begin{pmatrix} \frac{\partial \mathcal{J}}{\partial w_i} \\ \vdots \\ \frac{\partial \mathcal{J}}{\partial w_D} \end{pmatrix}$$

- ▣ This is the direction of fastest increase in  $\mathcal{J}$

- Update rule in vector form:

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} - \alpha \frac{\partial \mathcal{J}}{\partial \mathbf{w}} \\ &= \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)} \end{aligned}$$

- Hence, the gradient updates the weights in the direction of the fastest decrease.




## Gradient Descent: Visualization

### Gradient Descent: Visualization




For Linear Regression,  $J$  is bowl-shaped (“convex”)




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )






$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )



# Gradient Descent


- Why gradient descent, if we can find the optimum directly?
  - ▣ GD can be applied to a much broader set of models
  - ▣ GD can be easier to implement than direct solutions, especially with automatic differentiation software
  - ▣ For regression in high-dimensional spaces, GD is more efficient than direct solution (matrix inversion is an  $\mathcal{O}(D^3)$  algorithm).





## Feature Mappings

- Suppose we want to model the following data



- One option: fit a low degree polynomial; this is known as polynomial regression

$$y = w_3x^3 + w_2x^2 + w_1x + w_0$$

- Do we need to derive a whole new algorithm?



## Feature Mappings

- We get polynomial regression for free!
- Define the feature map

$$\phi(x) = \begin{pmatrix} 1 \\ x \\ x^2 \\ x^3 \end{pmatrix}$$


- Polynomial regression model:

$$y = \mathbf{w}^T \phi(x)$$


- All of the derivations and algorithms so far in the lecture remain the same!



# Fitting Polynomials




# Fitting Polynomials




# Fitting Polynomials

$$y = w_0 + w_1x + w_2x^2 + w_3x^3$$




# Fitting Polynomials

$$y = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_9x^9$$




## Generalization

- Underfitting: model is too simple — does not fit the data.





- Overfitting: model is too complex - fits perfectly, does not generalize.

$$y = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_9x^9$$



# Generalization

- Training and test error as a function of # training examples and # parameters:





# Logistic Regression



## Classification Problem

- The linear regression model assumes that the response variable  $y$  is quantitative (a value).
- In many situations, the response variable is instead qualitative (categorical).
  - ▣ Mango  $\in \{\text{Lengra, Harivanga, Amropoli}\}$
  - ▣ Email  $\in \{\text{Spam, Ham}\}$



## Classification Problem

- The process of estimating categorical outcomes using a set of features  $\mathbf{x}$  is called classification.
- Estimating a categorical response for an observation  $\mathbf{x}$  can be referred to as classifying that observation since it involves assigning the observation to a category, or class
- Often we are more interested in estimating the probabilities that  $\mathbf{x}$  belongs to each category in  $\mathcal{C}$
- The most probable category is then chosen as the class for the observation  $\mathbf{x}$




## Example: cat vs dog classification

Suppose that we measure the weight and height of some dogs and cats

- We want to learn the function  $f()$ . that can tell us if a given input vector  $\mathbf{x} = [x_1, x_2]^T$  is a dog or a cat

- ▣  $x_1$ : weight
- ▣  $x_2$ : height




**Quiz:** Which class would you classify the point marked as a yellow diamond?



# Logistic Regression

- **Purpose:** Estimate the probability that a set of input features  $\mathbf{x} \in \mathbb{R}^{d \times 1}$  belong to one of two classes  $y \in \{0, 1\}$



Define the linear combination quantity  $a = \sum_{i=0}^{d-1} x_i \cdot w_i = \mathbf{x}^T \cdot \mathbf{w}$

## $s(a)$ , Logistic Function Formula

$$s(a) = \frac{1}{1+e^{-a}}, \text{ if } a \gg 0, s(a) = 1 \text{ and if } a \ll 0, s(a) = 0.$$


# Logistic Regression

$$P(y = 1|\mathbf{x}) = s(a) = s(\mathbf{x}^T \cdot \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{x}^T \cdot \mathbf{w}}}$$

■ The output of  $s(\mathbf{x}^T \cdot \mathbf{w})$  is interpreted as a probability

▣  $\mathbf{x}^T \cdot \mathbf{w} \gg 0 \Rightarrow s(\mathbf{x}^T \cdot \mathbf{w}) \gg 0.5 \Rightarrow P(y = 1|\mathbf{x}) \approx 1$

▣  $\mathbf{x}^T \cdot \mathbf{w} \ll 0 \Rightarrow s(\mathbf{x}^T \cdot \mathbf{w}) \ll 0.5 \Rightarrow P(y = 1|\mathbf{x}) \approx 0$



## Logistic Regression Cost Function

- Suppose we have a dataset  $\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$  where  $\mathbf{x}^{(i)} \in \mathbb{R}^{d \times 1}$  and  $y^{(i)} \in \{0, 1\}, i = 1, \dots, N$ .
- Estimate a logistic regression model  $P(y^{(i)} = 1 | \mathbf{x}^{(i)}) = \frac{1}{1 + e^{-\mathbf{x}^{(i)T} \mathbf{w}}} \equiv \pi^{(i)}$
- The logistic regression cost function  $\mathcal{J}(\mathbf{w})$  is defined as:

$$\mathcal{J}(\mathbf{w}) = - \sum_{i=1}^N (y^{(i)} \cdot \log \pi^{(i)} + (1 - y^{(i)}) \cdot \log[1 - \pi^{(i)}])$$

where:

- $\pi^{(i)}$  is the probability of  $\mathbf{x}^{(i)}$  belonging to class 1.
- $\mathbf{x}^{(i)}$  is the input feature.
- $\mathbf{w}$  are the parameters to be learned.



## Quiz:

$$\mathcal{J}(\mathbf{w}) = - \sum_{i=1}^N (y^{(i)} \cdot \log \pi^{(i)} + (1 - y^{(i)}) \cdot \log[1 - \pi^{(i)}])$$

- 



## Logistic Regression Cost Function


Cost function interpretation

Suppose there is only one datum  $\mathcal{D} = \{(\mathbf{x}, y)\}$

$$\mathcal{J}(\mathbf{w}) = \begin{cases} -\log \pi, & \text{if } y = 1 \\ -\log[1 - \pi], & \text{if } y = 0 \end{cases}$$

Case  $y = 1$

$$\mathcal{J}(\mathbf{w}) = -\log \pi \quad \begin{array}{l} \mathcal{J}(\mathbf{w}) \approx 0 \text{ if } y = 1 \text{ and } \pi \approx 1 \\ \mathcal{J}(\mathbf{w}) \approx +\infty \text{ if } y = 1 \text{ and } \pi \approx 0 \end{array}$$



## Logistic Regression Cost Function


Cost function interpretation

Suppose there is only one datum  $\mathcal{D} = \{(\mathbf{x}, y)\}$

$$\mathcal{J}(\mathbf{w}) = \begin{cases} -\log \pi, & \text{if } y = 1 \\ -\log[1 - \pi], & \text{if } y = 0 \end{cases}$$

Case  $y = 0$

$\mathcal{J}(\mathbf{w}) = -\log[1 - \pi]$	$\mathcal{J}(\mathbf{w}) \approx 0$ if $y = 0$ and $\pi \approx 0$
	$\mathcal{J}(\mathbf{w}) \approx +\infty$ if $y = 0$ and $\pi \approx 1$



# Gradient Descent

Gradient Descent is used to find the parameters that minimize the cost function.

$$w_j = w_j - \alpha \frac{\partial \mathcal{J}(\mathbf{w})}{\partial w_j} \quad (1)$$

where:

- $\alpha$  is the learning rate.
- $\frac{\partial \mathcal{J}(\mathbf{w})}{\partial w_j}$  is the partial derivative of the cost function with respect to parameter  $w_j$ .





**Thank You!**